

What Is Serverless?

What is Serverless?

Serverless doesn't mean there are no servers, it means you don't care about them.

Serverless can be group into two areas:

- **Backend as a Service (BaaS)** Replacing server-side, self-managed components with off-the-shelf services
- **Functions as a Service (FaaS)** A new way of building and deploying server-side software, oriented around deploying individual functions

The key is that with both, you don't have to manage your own server hosts or server processes and can focus on business value!

What is Serverless?

Key characteristics

A Serverless service ...

- ... does not require managing a long-lived host or application instance
- ... self auto-scales and auto-provisions, dependent on load
- ... has implicit high availability
- ... has performance capabilities defined in terms other than host size/count
- ... has costs that are based on precise usage, up from and down to zero usage

Why Serverless?

- Shorter lead time
 - Reduced packaging and deployment complexity
- Increased flexibility of scaling
- Reduced resource cost
- Reduced labor cost
- Reduced risk

Drawbacks / Limitations of Serverless

Part 1/2

- Unpredictable costs
- Spinning up machines takes time - from a few seconds to minutes
- Most Serverless applications are stateless and the management of state can be somewhat tricky
- Higher latency due to inter-component communication over HTTP APIs and “cold starts”
- Problematic with downstream systems that cannot increase their capacity quickly enough
- Typically limited in how long each invocation is allowed to run
- Multitenancy problems

Drawbacks / Limitations of Serverless

Part 2/2

- Debugging is more complicated (a single request can travel between several machines and some of those machines disappear at times)
- Added work is needed to provide tracing and monitoring solutions, which can add complexity and cost to the project
- Security can be more demanding in a serverless environment
- Loss of control over
 - absolute configuration
 - the performance of Serverless components
 - issue resolution
 - security
- The difficulty of local testing
- Vendor lock-in unless you are using OSS projects like e.g. Knative

What is Knative?

An open-source community project which adds components for deploying, running, and managing applications on any Kubernetes in a Serverless way.

Two primary components:

- **Serving:** Supports deploying and serving of serverless applications and functions
- **Eventing:** Enables developers to use an event-driven architecture with serverless applications




Summary

The primary drivers for the adoption of Serverless are:



Developer
productivity



Platform elasticity



Cost savings


Going Serverless

with your Spring Boot applications

Demo

Running a Spring Boot application on Knative

Unleash the full potential of Serverless for our application



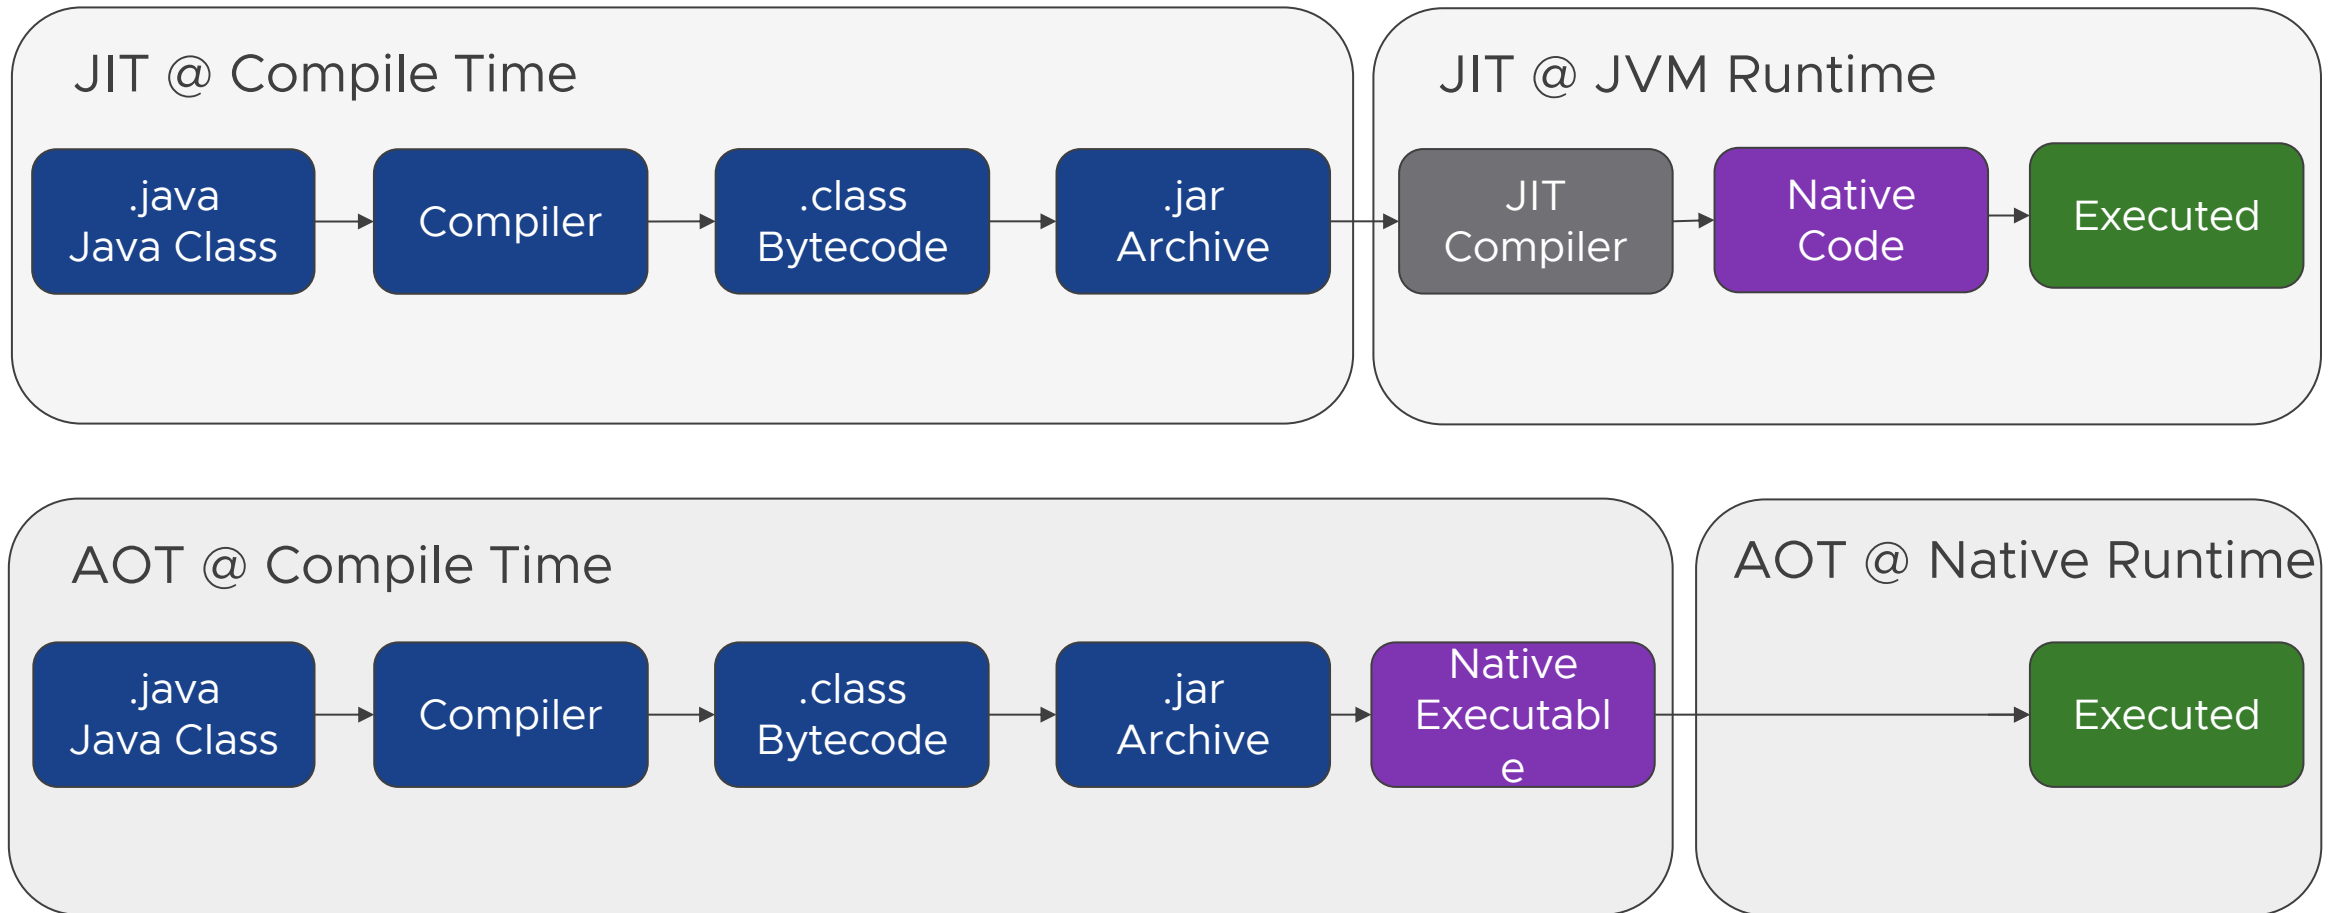
Faster startup time

Lower resource
consumption
(memory, CPU)

More cost
savings

Just-in-Time vs Ahead-of-Time

Interpretation vs Compilation



What are Native Images?

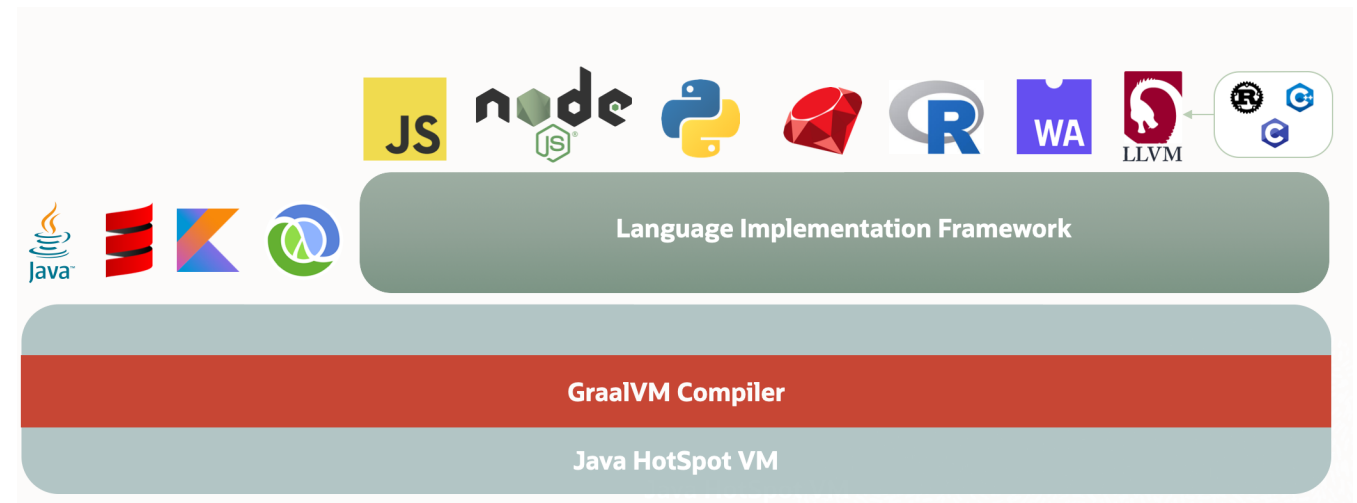
Smaller, faster and lower resource consumption

- Standalone executable of ahead-of-time compiled Java code
- Includes the application classes, classes from its dependencies, runtime library classes, and statically linked native code from JDK
- Runs without the need of a JVM, necessary components are included in a runtime system, called “Substrate VM”
- Specific to the OS and machine architecture for which it was compiled
- Requires fewer resources than regular Java applications running on a JVM


What is GraalVM?

A high-performance JDK distribution

- Basic idea: One VM that can execute applications written in Java and other JVM languages while also providing runtimes for JavaScript, Ruby, Python, and a number of other popular languages
- Started by Oracle: <https://graalvm.org/>
- GraalVM's polyglot capabilities make it possible to mix multiple programming languages in a single application while eliminating any foreign language call costs
- Still in “early adopter” mode, but matures quickly



Tradeoffs between JVM and Native Images



Key differences between JVM and GraalVM native image platform

- A static analysis of your application from the main entry point is performed at build time
- The unused parts are removed at build time
- Configuration is required for reflection, resources, and dynamic proxies
- Classpath is fixed at build time
- No class lazy loading: everything shipped in the executables will be loaded in memory on startup
- Some code will run at build time
- There are some [limitations](#) around some aspects of Java applications that are not fully supported



Spring Native provides incubating support for compiling Spring applications to lightweight native executables using the GraalVM native-image compiler.

The goal is to support compilation of existing or new Spring Boot applications to native executables.

Unchanged.

Get started with Spring Native

Create a new project on start.spring.io or use an existing one

The screenshot shows the Spring Initializr web form. It has three main sections: Project, Language, and Dependencies. The Project section includes options for Maven Project, Gradle Project, and Project Metadata (Group, Artifact, Name, Description, Package name). The Language section includes options for Java, Kotlin, and Groovy. The Dependencies section includes a button to add dependencies and a section for Spring Native [Experimental].

Annotations with blue arrows point to the following elements:

- The **2.5.4** radio button under the **Spring Boot** section.
- The **11** radio button under the **Java** section.
- The **ADD DEPENDENCIES...** button in the **Dependencies** section.

At the bottom of the form, there are three buttons: **GENERATE**, **EXPLORE**, and **SHARE...**.

Choose **Spring Boot 2.5.4** as your project's parent for the Spring Native 0.10.3 release

Java 8 and Java 11 supported for Native Images at this time in the GraalVM

Get started with Spring Native

Create a new project on start.spring.io or use an existing one

```
<parent>
  <groupId>org.springframework.boot</groupId>
  <artifactId>spring-boot-starter-parent</artifactId>
  <version>2.5.4</version>
  <relativePath/> <!-- lookup parent from repository -->
</parent>
```

```
<dependencies>
  <dependency>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter</artifactId>
  </dependency>
  <dependency>
    <groupId>org.springframework.experimental</groupId>
    <artifactId>spring-native</artifactId>
    <version>0.10.3</version>
  </dependency>
</dependencies>
```

```
<build>
  <plugins>
    <plugin>
      <groupId>org.springframework.boot</groupId>
      <artifactId>spring-boot-maven-plugin</artifactId>
      <configuration>
        <image>
          <builder>paketobuildpacks/builder:tiny</builder>
          <env>
            <BP_NATIVE_IMAGE>true</BP_NATIVE_IMAGE>
          </env>
        </image>
      </configuration>
    </plugin>
    <plugin>
      <groupId>org.springframework.experimental</groupId>
      <artifactId>spring-aot-maven-plugin</artifactId>
      <version>0.10.3</version>
    </plugin>
  </plugins>
</build>
```

- Use [Spring Boot 2.5.4](#) as your project's parent for the latest Spring Native release
- Add a dependency to the latest Spring Native library
e.g. [spring-native](#) version [0.10.3](#)
- Leverage Paketo Java Native Image Buildpack
- Configure the Spring Boot Maven Plugin to build a native image
- Add the Spring AOT plugin

Get started with Spring Native

Create a new project on start.spring.io or use an existing one

```
<profiles>
<profile>
  <id>native</id>
  <build>
    <plugins>
      <plugin>
        <groupId>org.graalvm.buildtools</groupId>
        <artifactId>native-maven-plugin</artifactId>
        <version>0.9.3</version>
        <executions>
          <execution>
            <id>test-native</id>
            <phase>test</phase>
            <goals>
              <goal>test</goal>
            </goals>
          </execution>
          <execution>
            <id>build-native</id>
            <phase>package</phase>
            <goals>
              <goal>build</goal>
            </goals>
          </execution>
        </executions>
      </plugin>
    </plugins>
  </build>
</profile>
</profiles>
```

■ Add the native build tools plugin

Two options to build a native image with Spring Native

Via Cloud-Native-Buildpacks

- Configure your build to use the Paketo Buildpacks
- Tell the buildpack to produce a native image
- The result is a small container image with the compiled native executable inside
- No local GraalVM installation needed
- Super easy to use
- Run `mvn spring-boot:build-image` to create a container with your application

Via GraalVM native image Maven plugin

- Configure your build to compile to a native executable
- Produces a native executable for the platform you are running on
- Requires GraalVM locally installed
- Also super easy to use
- Run `mvn -Pnative -DskipTests package` to build the application.

Demo

Running a Spring Boot application as Native Image on Knative

Performance expectations

Sample	On the JVM	Native application
spring-boot-hello-world	Build: 19 s Memory(RSS): 225M Startup time: 2.367s	Build: 04:21 min +1374% Memory(RSS): 68M -70% Startup time: 0.119s -95%
actuator-r2dbc-webflux	Build: 8s Memory(RSS): 640M Startup time: 3.0s	Build: 141s +1700% Memory(RSS): 86M -87% Startup time: 0.094s -97%
petclinic-jdbc	Build: 9s Memory(RSS): 417M Startup time: 2.6s	Build: 194s +2050% Memory(RSS): 101M -75% Startup time: 0.158s -94%

Can I keep the great Spring UX when going Native?

Use JVM locally, delegate Native Builds to CI/CD...

- Delegate Native Image building to the CI/CD pipeline, leveraging Spring Native and Paketo Buildpacks
- Unchanged Developer UX in your IDE, with the same JVM

The road ahead

Spring Boot 3, based on Spring Framework 6, is expected to provide first-class support for native application deployment, as well as an optimized footprint on the JVM.